KRISHNA SHARMA

krishnajan2004@gmail.com — 9315268810 — Portfolio — LinkedIn — GitHub

Work Experience

- **AI/ML Intern** — Zeronsec, Vadodara *(May '24 – July '24)*
  - Built a data pipeline processing 15 GB/day of cybersecurity logs using schema validation and streaming.
  - Improved throughput by 60% via bottleneck analysis and optimized parallel ETL workflows.
  - Automated data preprocessing reducing manual effort by 45% and designed scale-up plan for 50+ GB/day.

Projects

- **ShipStream — Full-Stack Deployment Platform** — Link
  - Built a Vercel-like deployment platform that supports GitHub OAuth, auto-builds, and live domain routing using Next.js 15 (App Router) and Node.js with a responsive admin dashboard (TailwindCSS, Shadcn/ui, Framer Motion).
  - Designed microservice architecture with Redis queues that manage parallel build workflows and persistent deployment history.
  - Integrated Cloudflare R2, Upstash Redis, and Render to achieve cost-free scalability and zero-downtime builds with real-time status tracking.
- **Sevak — AI-Powered Legal Chatbot Platform** — Link
  - Developed an AI chatbot using FastAPI and Gemini to explain Indian laws, handling 1,000+ user queries with 95% positive feedback.
  - Added OTP-based verification and role-based access to prevent spam and ensure verified submissions.
  - Built a ticket system with admin escalation and automated weekly data cleanup using APScheduler.
  - Integrated a location-based emergency contact feature for 50+ city zones, responding within 3 seconds.
  - Deployed via Docker with a React.js frontend and MongoDB backend.
- **Cross-Source Log Anomaly Detection System** — Link
  - Designed a distributed log analysis system using PySpark and parallel processing to handle 16+ heterogeneous sources (Windows, Linux, Hadoop, HDFS, Spark).
  - Implemented and benchmarked 9+ ML models for anomaly detection, including DANN-BERT, LoRA-BERT, and Hybrid-BERT, achieving scalable cross-domain generalization with hybrid embeddings and adversarial domain adaptation.
  - Optimized model training and data ingestion pipelines through multi-core parallelism, reducing total processing time by 65%.

Education

- **B-Tech CSE** — SRM University AP **CGPA: 7.85** (Oct '22 - Aug '26)
- **XII (ISC)** — St. Joseph's School (2022)

Technical Skills

Languages: JavaScript (ES6+), TypeScript, Python, Java, C++, SQL
Frontend: React.js, Next.js (App Router, Server Actions), TailwindCSS, Shadcn/ui, Recharts
Backend: Node.js, Express.js, Django, FastAPI, REST APIs, WebSocket, Prisma ORM, NextAuth.js
Database & Cache: PostgreSQL, MySQL, MongoDB, Redis (Upstash), Query Optimization
Machine Learning & Big Data: PySpark, TensorFlow, PyTorch, Scikit-learn, Hadoop, Data Pipelines
Cloud & DevOps: AWS (S3, EC2, SageMaker), Docker, CI/CD, Cloudflare, Vercel, Render
Tools: Git/GitHub, MLflow, Tableau, Power BI, Excel Analytics

Certifications

- AWS Certified Cloud Practitioner — Amazon Web Services — Verify
- Oracle Cloud Infrastructure 2025 Generative AI Professional — Oracle — Verify
- Oracle Cloud Infrastructure 2025 Certified Data Science Professional — Oracle — Verify